



# COMPUTER ORGANIZATION

## 2023 SPRING

### FINAL PROJECT PART 2

**0811128 林洋宏**

**310511050 張祐誠**

**109612029 郭子維**

## Function Correctness&Best Performance (80%)

Cachline size = 64B, cache size = 2kB (in transpose/64B/), matrix size = 136\*136

Baseline=404,972 $\pm$ 10,000

Our performance=393,716

```
root@5d5bd004ffe3:/home/gem5/transpose_student/64B#  
info: Increasing stack size by one page.  
info: Increasing stack size by one page.  
info: Increasing stack size by one page.  
info: Increasing stack size by one page.  
info: Increasing stack size by one page.  
info: Increasing stack size by one page.  
info: Increasing stack size by one page.  
info: Increasing stack size by one page.  
info: Increasing stack size by one page.  
info: Increasing stack size by one page.  
info: Increasing stack size by one page.  
info: Increasing stack size by one page.  
info: Increasing stack size by one page.  
info: Increasing stack size by one page.  
info: Increasing stack size by one page.  
info: Increasing stack size by one page.  
info: Increasing stack size by one page.  
info: Increasing stack size by one page.  
=====  
correct answer \^o^/  
=====  
Exiting @ tick 3268817000 because exiting with less active thread context  
-----  
system.cpu.numCycles          393716 # number of cpu cycles simulated  
system.cpu.dcache.overall_miss_rate::cpu.data    0.108757 # miss rate for overall accesses  
system.cpu.icache.overall_miss_rate::cpu.inst    0.001966 # miss rate for overall accesses  
system.cpu.dcache.demand_hits::cpu.data         48840 # number of demand (read+write) hits  
system.cpu.dcache.demand_misses::cpu.data       5937 # number of demand (read+write) misses  
-----  
root@5d5bd004ffe3:/home/gem5/transpose_student/64B#
```

Cachline size = 128B, cache size = 4kB (in transpose/128B/), matrix size = 144\*144

Baseline=409,695 $\pm$ 10,000

Our performance=313,183

```
root@5d5bd004ffe3:/home/gem5/transpose_student/128B#  
info: Increasing stack size by one page.  
info: Increasing stack size by one page.  
info: Increasing stack size by one page.  
info: Increasing stack size by one page.  
info: Increasing stack size by one page.  
info: Increasing stack size by one page.  
info: Increasing stack size by one page.  
info: Increasing stack size by one page.  
info: Increasing stack size by one page.  
info: Increasing stack size by one page.  
info: Increasing stack size by one page.  
info: Increasing stack size by one page.  
info: Increasing stack size by one page.  
info: Increasing stack size by one page.  
info: Increasing stack size by one page.  
info: Increasing stack size by one page.  
=====  
correct answer \^o^/  
=====  
Exiting @ tick 3530160000 because exiting with last active thread context  
  
-----  
system.cpu.numCycles          313183      # number of cpu cycles simulated  
system.cpu.dcache.overall_miss_rate::cpu.data    0.140356      # miss rate for overall accesses  
system.cpu.icache.overall_miss_rate::cpu.inst     0.000577      # miss rate for overall accesses  
system.cpu.dcache.demand_hits::cpu.data           39872        # number of demand (read+write) hits  
system.cpu.dcache.demand_misses::cpu.data         6510         # number of demand (read+write) misses  
-----  
root@5d5bd004ffe3:/home/gem5/transpose_student/128B# vim todo_128.cpp  
root@5d5bd004ffe3:/home/gem5/transpose_student/128B# vim todo_128.cpp  
root@5d5bd004ffe3:/home/gem5/transpose_student/128B#
```

## Report (20%)

### o Screenshot of your results

The screenshot shows the Docker Desktop interface. The left sidebar contains navigation options: Containers, Images, Volumes, Dev Environments (BETA), Extensions, and Add Extensions. The main area displays a list of containers under the heading 'Containers'. A search bar and a toggle for 'Only show running containers' are at the top. The container list has columns for Name, Image, Status, Port(s), Last started, and Actions. Two containers are listed: '2023CO\_FP' (Running) and 'kind\_rosalind' (Exited). The '2023CO\_FP' container is highlighted with a red box. The 'Last started' column is also highlighted with a red box, showing '3 seconds ago' for the '2023CO\_FP' container.

|                          | Name                                 | Image                                     | Status       | Port(s) | Last started  | Actions |
|--------------------------|--------------------------------------|---|--------------|---------|---------------|---------|
| <input type="checkbox"/> | <b>2023CO_FP</b><br>5d5bd004ffe3     | <a href="#">cksc33milktea/gem5_2022co</a> | Running      |         | 3 seconds ago |         |
| <input type="checkbox"/> | <b>kind_rosalind</b><br>eb2226ac626c | <a href="#">docker/getting-started</a>    | Exited (255) | 80:80   | 24 days ago   |         |

Showing 2 items

RAM 3.55 GB CPU 8.08% Not connected to Hub v4.19.0

1. Start the container and enter
2. its command line (Open Docker first):  
`docker start 2023CO_FP`

docker cp YOUR\_DOWNLOAD\_PATH\transpose 2023CO\_FP:/home/gem5/

```
命令提示字元
Microsoft Windows [版本 10.0.22621.1702]
(c) Microsoft Corporation. 著作權所有，並保留一切權利。

C:\Users\s6106>docker cp C:\Users\s6106\Desktop\計算機組織\transpose_student\transpose 2023CO_FP:/home/gem5/
CreateFile C:\Users\s6106\Desktop\計算機組織\transpose_student\transpose: The system cannot find the file specified.

C:\Users\s6106>docker cp C:\Users\s6106\Desktop\計算機組織\transpose_student 2023CO_FP:/home/gem5/
Successfully copied 0B to 2023CO_FP:/home/gem5/
error during connect: this error may indicate that the docker daemon is not running: Put "http://%2F%2F.%2Fpipe%2Fdocker_engine/v1.24/containers/2023CO_FP/archive?noOverwriteDirNonDir=true&path=%2Fhome": open //./pipe/docker_engine: The system cannot find the file specified.

C:\Users\s6106>docker start 2023CO_FP
2023CO_FP

C:\Users\s6106>docker exec -ti 2023CO_FP bash
root@5d5bd004ffe3:/# docker cp C:\Users\s6106\Desktop\計算機組織\transpose_student 2023CO_FP:/home/gem5/
bash: docker: command not found
root@5d5bd004ffe3:/# exit
exit

C:\Users\s6106>docker cp C:\Users\s6106\Desktop\計算機組織\transpose_student 2023CO_FP:/home/gem5/
Successfully copied 13.8kB to 2023CO_FP:/home/gem5/

C:\Users\s6106>
```

3. Start docker image:

```
docker start 2023CO_FP
```

```
docker exec -ti 2023CO_FP bash
```

4. Revise todo\_128.cpp and todo\_64.cpp:

```
cd home/gem5/transpose_student/
```

```
cd 128B/ or cd 64B/
```

```
root@5d5bd004ffe3:/home/gem5/transpose_student/64B
../build/X86/gem5.opt ../configs/example/se.py --cpu-type=Minicpu --l1d_size=2kB --l1d_assoc=4 --cacheline_size=64 --l1i_size=4kB --l1i_assoc=2 --caches --mem-type=Ramulator --ramulator-config=../configs/ramulator/DDR4-config.cfg -c main
echo "-----";
cat m5out/stats.txt |grep "numCycles" | head -n 1;
cat m5out/stats.txt |grep "dcache.overall_miss_rate" | head -n 1;
cat m5out/stats.txt |grep "icache.overall_miss_rate" | head -n 1;
cat m5out/stats.txt |grep "dcache.demand_hits" | head -n 1;
cat m5out/stats.txt |grep "dcache.demand_misses" | head -n 1;
echo "-----";
^
^
^
^
```


Run run\_sim.sh:

```
make
```


```
./run_sim.sh
```

Blocking method:

| 128B     |              | 64B      |              |
|----------|--------------|----------|--------------|
| tileSize | cycles/ticks | tileSize | cycles/ticks |
| 1        | 898685       | 1        | 819856       |
| 2        | 560275       | 2        | 522050       |
| 4        | 611717       | 4        | 623252       |
| 8        | 575781       | 8        | 643762       |
| 16       | 557743       | 16       | 582102       |
| 18       | 547149       | 18       | 561672       |
| 20       | 544869       | 20       | 548902       |
| 22       | 544495       | 22       | 564570       |
| 24       | 560429       | 24       | 580778       |
| 32       | 592487       | 32       | 631718       |
| 64       | 859655       | 64       | 801758       |



The best case for 128B cahceline with blocking is 544495 cycles.




The best case for 64B cahceline with blocking is 522050 cycles.

Conclusion: optimizing using only blocking method is not enough.

## Blocking + prefetch:

| 128B     |              | 64B      |              |
|----------|--------------|----------|--------------|
| tileSize | cycles/ticks | tileSize | cycles/ticks |
|          |              |          |              |
| 2        | 611599       | 2        | 555350       |
| 4        | 514629       | 4        | 505410       |
| 8        | 464591       | 8        | 531516       |
| 10       | 454189       | 10       | 488874       |
| 12       | 425307       | 12       | 439280       |

|     |        |     |        |
|-----|--------|-----|--------|
| 14  | 437547 | 14  | 464768 |
| 16  | 448929 | 16  | 479504 |
| 18  | 468801 | 18  | 466308 |
| 20  | 499173 | 20  | 462974 |
| 22  | 529925 | 22  | 489492 |
| 24  | 570107 | 24  | 500724 |
| 32  | 565573 | 32  | 542104 |
| 34  | 591139 | 34  | 544908 |
| 64  | 602315 | 64  | 560962 |
| 68  | 612863 | 68  | 569224 |
| 256 | 612855 | 256 | 557474 |



The best case for blocking + prefetch for 64B cachline is 439280 cycles.


Conclusion: we are on the right track!

## Blocking + unroll:


|          |              |        |          |              |        |
|----------|--------------|--------|----------|--------------|--------|
| 128B     |              |        | 64B      |              |        |
| tileSize | unrollFactor | cycle  | tileSize | unrollFactor | cycle  |
| 64       | 8            | 438636 | 64       | 8            | 484718 |
| 128      | 8            | 427148 | 128      | 8            | 475938 |
| 144      | 8            | 452804 | 136      | 8            | 489164 |
| 256      | 8            | 437432 | 256      | 8            | 474514 |
| 1024     | 8            | 437432 | 1024     | 8            | 474514 |
| 2048     | 8            | 437432 | 2048     | 8            | 474514 |
| 64       | 12           | 452932 | 64       | 12           | 535552 |
| 128      | 12           | 427266 | 128      | 12           | 498440 |
| 143      | 12           | 425227 | 136      | 12           | 423650 |
| 144      | 12           | 317809 | 137      | 12           | 414534 |




|      |    |        |      |    |               |
|------|----|--------|------|----|---------------|
| 145  | 12 | 313183 | 256  | 12 | <b>414534</b> |
| 256  | 12 | 313183 | 1024 | 12 | <b>414534</b> |
| 1024 | 12 | 313183 | 2048 | 12 | <b>414534</b> |
| 2048 | 12 | 313183 |      |    |               |




Let unrollfactor(upper) be set to 8, and tileSize(lower) be set to 64~2048, we can see it converges to 437432 cycles.



Let unrollfactor(upper) be set to 12, and tileSize(lower) be set to 64~2048, we can see it converges to **313183 cycles**.(great improvement)



Let unrollfactor(upper) be set to 8, and tileSize(lower) be set to 64~2048, we can see it converges to 474514 cycles.




Let unrollfactor(upper) be set to 12, and tileSize(lower) be set to 64~2048, we can see it converges to 414534 cycles.

Conclusion, we have reached the goal! But there are still possibilities of optimizing even further on 64B cache line.


### Blocking + Unrolling + loop interchange+ prefetch:

| 128B     |              |        | 64B      |              |        |
|----------|--------------|--------|----------|--------------|--------|
| tileSize | unrollFactor | cycle  | tileSize | unrollFactor | cycle  |
| 64       | 8            | 417462 | 64       | 8            | 474054 |
| 128      | 8            | 410028 | 128      | 8            | 464948 |
| 144      | 8            | 446691 | 136      | 8            | 463442 |
| 256      | 8            | 446351 | 256      | 8            | 463944 |
| 1024     | 8            | 446351 | 1024     | 8            | 463944 |


|      |    |        |      |    |        |
|------|----|--------|------|----|--------|
| 2048 | 8  | 446351 | 2048 | 8  | 463944 |
| 64   | 12 | 447460 | 64   | 12 | 482296 |
| 128  | 12 | 419088 | 128  | 12 | 431100 |
| 144  | 12 | 328458 | 136  | 12 | 393716 |
| 256  | 12 | 328458 | 256  | 12 | 393716 |
| 1024 | 12 | 328458 | 1024 | 12 | 393716 |
| 2048 | 12 | 328458 | 2048 | 12 | 393716 |




Let unrollfactor(upper) be set to 8, and tileSize(lower) be set to 64~2048, we can see it converges to 446351 cycles.



Let unrollfactor(upper) be set to 12, and tileSize(lower) be set to 64~2048, we can see it converges to **328458 cycles**.



Let unrollfactor(upper) be set to 8, and tileSize(lower) be set to 64~2048, we can see it converges to 463944 cycles.



Let unrollfactor(upper) be set to 8, and tileSize(lower) be set to 64~2048, we can see it converges to **393716 cycles.**

**Conclusion, 64B cacheline has been further optimized.**

◦ Explain how did you optimize your code according to two different caches

## 128B: Best optimization → Blocking + unrolling

```
void transpose(const uint64_t a[M][M], uint64_t b[M][M]){
    const int tileSize = 1024; // Size of the tiles
```



```

const int unrollFactor = 8; // Number of loop iterations to process per
iteration
for (int ii = 0; ii < M; ii += tileSize) {
    for (int jj = 0; jj < M; jj += tileSize) {
        for (int i = ii; i < min(ii + tileSize, M); i += unrollFactor) {
            for (int j = jj; j < min(jj + tileSize, M); j++){
                for (int k = 0; k < unrollFactor; k++) {
                    if (i + k < M) {
                        b[j][i + k] = a[i + k][j];
                    }
                }
            }
        }
    }
}

```

In this code, the matrix transposition is divided into smaller tiles of size 'tileSize', and the outer two loops iterate over these tiles. However, there is a difference in the inner loops. The loop iteration is unrolled for 'unrollFactor' times in the 'i' loop, and the 'j' loop is not unrolled. The 'k' loop iterates over the unrolled iterations.

Here's how these optimizations are associated with the cache characteristics:

1. **Cache Line Size:** The cache line size is 128 bytes. By transposing the matrix in tiles of size tileSize (1024), the code ensures that each tile fits within a few cache lines (1024 bytes / 128 bytes = 8 cache lines). This alignment with the cache line size improves cache utilization and reduces cache misses.
2. **Cache Size:** The cache size is 4kB, which is larger than the tile size (1024). This allows for efficient utilization of the cache. Each tile can fit within the cache, ensuring that a significant portion of the accessed data resides in cache memory. This reduces cache misses and improves overall performance.
3. **LRU Policy:** The LRU policy determines which cache lines to evict when the cache is full. By using loop tiling, the code enhances spatial locality and increases the chances of reusing cache lines within each tile. This aligns with

the LRU policy by maximizing cache hits within each set and reducing cache evictions.

4. Associativity: The cache has an associativity of 4, meaning each cache set can hold up to 4 cache lines. With loop tiling, the code accesses data within each tile in a localized manner, increasing the chances of cache hits within each set. This reduces conflicts and improves cache utilization, aligning with the given associativity.

In summary, the revised code optimizes cache utilization by employing **loop tiling and unrolling techniques**. The tiles fit well within the cache lines, ensuring efficient utilization of the cache. The cache size of 4kB provides large enough space for the tiles, reducing cache misses. The LRU policy and associativity of 4 align well with the code's focus on spatial locality, improving cache hit rates and overall performance. Therefore, this method can be optimized for a cache with a cache line size of 128B, cache size of 4kB, LRU policy, and associativity of 4 by effectively utilizing the cache through loop tiling and unrolling, reducing cache misses, and maximizing cache hit rates.

## 64B: Best optimization → Blocking + unrolling + loop

### interchange + prefetch

```
#include <immintrin.h> // Required for prefetching

void transpose(const uint64_t a[M][M], uint64_t b[M][M]) {
    const int tileSize = 1024; // Size of the tiles
    const int unrollFactor = 8; // Number of loop iterations to process
    per iteration

    for (int ii = 0; ii < M; ii += tileSize) {
        for (int jj = 0; jj < M; jj += tileSize) {
            for (int j = jj; j < min(jj + tileSize, M); j += unrollFactor)
            {
                // Prefetch the next tile of matrix a
                _mm_prefetch((char*)&a[ii + tileSize][j], _MM_HINT_T0);

                for (int i = ii; i < min(ii + tileSize, M); i++) {
```

```

        for (int k = 0; k < unrollFactor; k++) {
            if (j + k < M) {
                b[j + k][i] = a[i][j + k];
            }
        }
    }

    // Prefetch the next set of elements from matrix a
    _mm_prefetch((char*)&a[ii][j + unrollFactor], _MM_HINT_T0);
}
}
}
}

```

In this code, the matrix transposition is divided into smaller tiles of size 'tileSize'. The outer two loops iterate over these tiles, and the inner loops perform the actual transposition within each tile. By processing smaller tiles at a time, the code takes advantage of spatial locality and reduces the number of cache misses.

The prefetching instructions (`_mm_prefetch`) are used to explicitly request the cache system to bring data into the cache ahead of time. This helps to minimize the impact of cache misses by **fetching data from memory in advance, anticipating future usage.**

The Loop Interchange instructions: in the innermost loop, we were accessing elements of array `a` in a column-major manner (`a[i + k][j]`). By interchanging the loop order, you can access the elements **in a row-major manner (`a[j][i + k]`)**, which can improve cache locality and potentially lead to better performance.

Here's how these optimizations are associated with the cache characteristics:

1. **Cache Line Size:** The cache line size is 64 bytes. By transposing the matrix in tiles of size `tileSize`, which is 1024, the code ensures that each tile fits within a few cache lines (1024 bytes / 64 bytes = 16 cache lines). This improves cache utilization as the data accessed within **each tile is likely to reside in the same cache lines, reducing cache misses.**
2. **Cache Size:** The cache size is 2kB. The tile size of 1024 allows for efficient utilization of the cache. Each tile is smaller than the cache size, ensuring that a significant portion of the accessed data resides within the cache. This

reduces the number of cache misses and improves overall performance.

3. LRU Policy: The LRU policy determines which cache lines to evict when the cache is full. By using loop tiling, the code enhances spatial locality and increases the chances of reusing cache lines within each tile. This reduces cache evictions and improves cache hit rates, thereby benefiting from the LRU policy.
4. Associativity: The cache has an associativity of 4, meaning each cache set can hold up to 4 cache lines. With loop tiling, the code accesses data within each tile in a localized manner, increasing the chances of cache hits within each set. This reduces conflicts and improves cache utilization.
- 5.

In summary, by employing **loop tiling and prefetching techniques**, the code optimizes cache utilization by exploiting **spatial locality**. It ensures that the accessed data within each tile fits within cache lines, minimizes cache misses, and maximizes cache hit rates. This leads to improved performance, especially when the cache has a cache line size of 64B, cache size of 2kB, LRU policy, and associativity of 4.

○ Compare the result (cycles, miss rate of data cache) between the optimized and no-optimized codes

## Simplest implementation in todo\_64.cpp:

root@5d5bd004ffe3: /home/gem5/transpose\_student/64B

```
// Example program
#include <iostream>
#include <string>
#include <stdlib.h>
#include <stdint.h>
#include <bits/stdc++.h>
#define M 136

using namespace std;

void transpose(const uint64_t a[M][M], uint64_t b[M][M]){
    /*=====only modify in this region=====
    for (int i = 0; i < M; i++) {
        for (int j = 0; j < M; j++) {
            b[j][i] = a[i][j];
        }
    }
    */
}
```

todo\_64.cpp" [dos] 20L, 501C

16,3-24

All

root@5d5bd004ffe3: /home/gem5/transpose\_student/64B

```
info: Increasing stack size by one page.
info: Increasing stack size by one page.
info: Increasing stack size by one page.
info: Increasing stack size by one page.
info: Increasing stack size by one page.
info: Increasing stack size by one page.
info: Increasing stack size by one page.
info: Increasing stack size by one page.
info: Increasing stack size by one page.
info: Increasing stack size by one page.
info: Increasing stack size by one page.
info: Increasing stack size by one page.
info: Increasing stack size by one page.
info: Increasing stack size by one page.
info: Increasing stack size by one page.
```

correct answer \^o^/

**as expected, too many cycles, cycle is around 815274**

Exiting @ tick 3409444000 because exiting with last active thread context

```
-----
system.cpu.numCycles                815274                # number of cpu cycles simulated
system.cpu.dcache.overall_miss_rate::cpu.data    0.562470        # miss rate for overall accesses
system.cpu.icache.overall_miss_rate::cpu.inst    0.000015        # miss rate for overall accesses
system.cpu.dcache.demand_hits::cpu.data          16186           # number of demand (read+write) hits
system.cpu.dcache.demand_misses::cpu.data        20808           # number of demand (read+write) misses
-----
```

### Blocking + Unrolling + loop interchange+ prefetch:

```
root@5d5bd004ffe3:/home/gen5/transpose_student/64B#  
info: Increasing stack size by one page.  
info: Increasing stack size by one page.  
info: Increasing stack size by one page.  
info: Increasing stack size by one page.  
info: Increasing stack size by one page.  
info: Increasing stack size by one page.  
info: Increasing stack size by one page.  
info: Increasing stack size by one page.  
info: Increasing stack size by one page.  
info: Increasing stack size by one page.  
info: Increasing stack size by one page.  
info: Increasing stack size by one page.  
info: Increasing stack size by one page.  
info: Increasing stack size by one page.  
info: Increasing stack size by one page.  
info: Increasing stack size by one page.  
info: Increasing stack size by one page.  
  
=====
```

best case yet!

```
correct answer ^o^/  
  
=====
```


Exiting @ tick 3268817000 because exiting with last active thread context

```
=====
```

|   |          |  |
|---|----------|--|
| system.cpu.numCycles                          | 393716   | # number of cpu cycles simulated       |
| system.cpu.dcache.overall_miss_rate::cpu.data | 0.108757 | # miss rate for overall accesses       |
| system.cpu.dcache.overall_miss_rate::cpu.inst | 0.001966 | # miss rate for overall accesses       |
| system.cpu.dcache.demand_hits::cpu.data       | 48840    | # number of demand (read+write) hits   |
| system.cpu.dcache.demand_misses::cpu.data     | 5937     | # number of demand (read+write) misses |

```
=====
```

root@5d5bd004ffe3:/home/gen5/transpose\_student/64B#







## Blocking + unrolling


```
root@5d5bd004ffe3: /home/gem5/transpose_student/128B
```

```
info: Increasing stack size by one page.  
info: Increasing stack size by one page.  
info: Increasing stack size by one page.  
info: Increasing stack size by one page.  
info: Increasing stack size by one page.  
info: Increasing stack size by one page.  
info: Increasing stack size by one page.  
info: Increasing stack size by one page.  
info: Increasing stack size by one page.  
info: Increasing stack size by one page.  
info: Increasing stack size by one page.  
info: Increasing stack size by one page.  
info: Increasing stack size by one page.  
info: Increasing stack size by one page.  
info: Increasing stack size by one page.  
info: Increasing stack size by one page.  
=====  
correct answer ^o^/  
=====
```

Exiting @ tick 3530160000 because exiting with last active thread context

|   | 313183   | # number of cpu cycles simulated       |
|---|----------|--|
| system.cpu.numCycles                          |          |  |
| system.cpu.dcache.overall_miss_rate::cpu.data | 0.140356 | # miss rate for overall accesses       |
| system.cpu.icache.overall_miss_rate::cpu.inst | 0.000577 | # miss rate for overall accesses       |
| system.cpu.dcache.demand_hits::cpu.data       | 39872    | # number of demand (read+write) hits   |
| system.cpu.dcache.demand_misses::cpu.data     | 6510     | # number of demand (read+write) misses |

```
-----  
root@5d5bd004ffe3:/home/gem5/transpose_student/128B# vim todo_128.cpp  
root@5d5bd004ffe3:/home/gem5/transpose_student/128B# vim todo_128.cpp  
root@5d5bd004ffe3:/home/gem5/transpose_student/128B#
```



|                                    |               |           |
|------------------------------------|---------------|-----------|
|                                    | Non-optimized | optimized |
| Cycles                             | 897502        | 313183    |
| dcache.overall_miss_rate::cpu.data | 0.531237      | 0.140336  |
| icache.overall_miss_rate::cpu.inst | 0.000048      | 0.000577  |
| dcache.demand_hits::cpu.data       | 19441         | 39872     |
| dcache.demand_misses::cpu.data     | 22032         | 6510      |